

Lab 1: Matrix Multiplication

姜海天 19307110022

This report is to record the performance optimization of an matrix multiplication problem.

In this problem, we are going to calculate $\alpha \cdot AB + \beta \cdot C$, where $A \in \mathbb{R}^{m \times k}$, $B \in \mathbb{R}^{k \times n}$, $C \in \mathbb{R}^{m \times n}$, and store the answer in C .

0.1 Testing Environment

All the experiments are conducted on a server equipped with dual 64-core, 128-thread `AMD EPYC 7742 64-Core` CPU. The peak double-precision performance for a single CPU is 3.48 Tflops. L1d cache: 4MiB, L1i cache: 4MiB, L2 cache: 64MiB, L3 cache: 512MiB

The compiler used for the experiments is `clang++ 14.0.6`.

All the test uses `./executable 4096 4096 4096` as the test command.

0.2 Computation Amount

If use the algorithm stated in the naive implementation, there will be $mn + 3mnk$ floating-point operations. If $m = n = k = 4096$ in my test case, there is 206.2G floating-point operations. So the theoretical minimum computing time is $206.2G / (2 * 3.48T/s) = 30ms$.

0.3 Main Results

Optimization	Running Time(s)	Relative Speedup	Absolute Speedup	Percent of Peak
Naive	2246.3	1	1	0.0013%
+ interchange loops	224.76	10	10	0.013%
+ optimization flags	38.85	5.78	57.8	0.077%
+ reduced computation	32.36	/	69.4	0.062%
Parallel loops	1.40	4.2	1609	2.15%
+tiling	0.34	27.8	6691	8.94%
+compiler vectorization	0.32	1.04	6941	9.27%

1 Naive Implementation

We use the naive implementation without any compiler optimization as a baseline.

```
1 void student_gemm(int m, int n, int k,  
2                 const double * A, const double * B, double * C,  
3                 double alpha, double beta,  
4                 int lda, int ldb, int ldc) {  
5     for (int i = 0; i < m; i++) {  
6         for (int j = 0; j < n; j++) {  
7             C[i + j * ldc] *= beta;  
8             for (int p = 0; p < k; p++) {  
9                 C[i + j * ldc] += alpha * A[i + p * lda] * B[p + j * ldb];  
10            }  
11        }  
12    }  
13 }
```

The computing time for the naive method is **37min26s (2246252.1556ms)**.

Speedup: **1x**, percent of peak: **0.0013%**.

```
1 $ ~/opencilk/bin/clang++ gemm.cpp -o 1-naive  
2 $ ./1-naive-1 4096 4096 4096  
3 input: 4096 x 4096 x 4096  
4 minimal time spent: 2246252.1556 ms  
5 result: correct (err = 0.000000e+00)
```

The cache miss rate for the naive method: **4.7%** by the `valgrind` command.

2 Loop Order

Change the loop order from `i,j,p` in the naive implementation to `j,p,i` so that we have the locality for columns, since the matrices are stored in column major.

```
1 void student_gemm(int m, int n, int k,  
2                 const double * A, const double * B, double * C,  
3                 double alpha, double beta,  
4                 int lda, int ldb, int ldc) {  
5     // compute  $\beta \cdot C$  first  
6     for (int j = 0; j < n; ++j) {  
7         for (int i = 0; i < m; ++i) {  
8             C[i + j * ldc] *= beta;  
9         }  
10    }  
11    for (int j = 0; j < n; j++) {  
12        for (int p = 0; p < k; p++) {  
13            // put the whole column for C and A in cache  
14            for (int i = 0; i < m; i++) {  
15                C[i + j * ldc] += alpha * A[i + p * lda] * B[p + j * ldb];  
16            }  
17        }  
18    }  
19 }
```

The computing time for the better loop order is **3min44s (224764.6316ms)**.

Speedup: **10x**, percent of peak: **0.013%**.

```
1 $ ~/opencilk/bin/clang++ gemm.cpp -o 2-loop-order  
2 $ ./2-loop-order 4096 4096 4096  
3 input: 4096 x 4096 x 4096  
4 minimal time spent: 224764.6316 ms  
5 result: correct (err = 0.000000e+00)
```

The cache miss rate for the naive method: **1.1%** by the `valgrind` command.

3 Compiler Optimization

The most aggressive optimization `-O3` got the best performance, and the running time is only **38.8s** (**38847.4110ms**) now.

Speedup: **57.8x**, percent of peak: **0.077%**.

```
1 $ ~/opencilk/bin/clang++ gemm.cpp -o 3-01 -O1
2 $ ./3-01 4096 4096 4096
3 input: 4096 x 4096 x 4096
4 minimal time spent: 44184.6329 ms
5 result: correct (err = 0.000000e+00)
6
7 $ ~/opencilk/bin/clang++ gemm.cpp -o 3-02 -O2
8 $ ./3-02 4096 4096 4096
9 input: 4096 x 4096 x 4096
10 minimal time spent: 39349.4487 ms
11 result: correct (err = 0.000000e+00)
12
13 $ ~/opencilk/bin/clang++ gemm.cpp -o 3-03 -O3
14 $ ./3-03 4096 4096 4096
15 input: 4096 x 4096 x 4096
16 minimal time spent: 38847.4110 ms
17 result: correct (err = 0.000000e+00)
```

4 Reduce Floating-point Operations

The current implementation calculates

$$c_{i,j} = \beta \cdot c_{i,j} + \sum_p \alpha \cdot a_{i,p} \cdot b_{p,j}$$

There are lots of unnecessary floating-point operations. The calculation with reduced floating-point operation is

$$c_{i,j} = \alpha \cdot \left(\frac{\beta}{\alpha} \cdot c_{i,j} + \sum_p a_{i,p} \cdot b_{p,j} \right)$$

Now the amount of floating-point operations have been reduced from $mn + 3mnk$ to $2mn + 2mnk$. However, it requires more read and write of memory. So there is trade-off.

This method will not be suitable for the further optimization aiming at reducing memory operations because it brings more memory operations naturally.

Implementation:

```

1 void student_gemm(int m, int n, int k,
2                 const double * A, const double * B, double * C,
3                 double alpha, double beta,
4                 int lda, int ldb, int ldc) {
5     double ratio = beta / alpha;
6     //  $C \leftarrow (\beta/\alpha)C$ 
7     for (int j = 0; j < n; ++j) {
8         for (int i = 0; i < m; ++i) {
9             C[i + j * ldc] *= ratio;
10        }
11    }
12    //  $C \leftarrow C + AB$ 
13    for (int j = 0; j < n; j++) {
14        for (int p = 0; p < k; p++) {
15            for (int i = 0; i < m; i++) {
16                C[i + j * ldc] += A[i + p * lda] * B[p + j * ldb];
17            }
18        }
19    }
20    //  $C \leftarrow \alpha C$ 
21    for (int j = 0; j < n; ++j) {
22        for (int i = 0; i < m; ++i) {
23            C[i + j * ldc] *= alpha;
24        }
25    }
26 }

```

Experiment: The running time is **32.36s**.

Speedup: **69.4x**, percent of peak: **0.062%**. The percent of peak dropped because the amount of floating-point operations is reduced and more memory access is required.

```

1 $ ~/opencilk/bin/clang++ gemm.cpp -o 4-coeff-03 -O3
2 $ ./4-coeff-03 4096 4096 4096
3 input: 4096 x 4096 x 4096
4 minimal time spent: 32360.3723 ms
5 result: correct (err = 2.913225e-13)

```

5 Multi-Core Parallel Computation

Add parallelable for-loop for the reordered loop version. The version with excessive memory access doesn't perform better because there are too much write to memory, which is much slower than floating-point operations.

```
1 void student_gemm(int m, int n, int k,
2                 const double * A, const double * B, double * C,
3                 double alpha, double beta,
4                 int lda, int ldb, int ldc) {
5     cilk_for (int j = 0; j < n; ++j) {
6         for (int i = 0; i < m; ++i) {
7             C[i + j * ldc] *= beta;
8         }
9     }
10    cilk_for (int j = 0; j < n; j++) {
11        for (int p = 0; p < k; p++) {
12            for (int i = 0; i < m; i++) {
13                C[i + j * ldc] += alpha * A[i + p * lda] * B[p + j * ldb];
14            }
15        }
16    }
17 }
```

Experiment: The running time is 1.4s.

Speedup: **1609x**, percent of peak: **2.15%**.

```
1 $ ~/opencilk/bin/clang++ gemm.cpp -fopencilk -O3 -o 5-cilk-lo
2 $ ./5-cilk-lo 4096 4096 4096
3 input: 4096 x 4096 x 4096
4 minimal time spent: 1395.9282 ms
5 result: correct (err = 0.000000e+00)
```

6 Tiling

I added the tiling size `s` as an optional command line parameter. The default one is 128, which is the empirical best size.

Implementation:

```
1 // ...
2 int s = 128; // global default value for s
3 void student_gemm(int m, int n, int k,
4                  const double * A, const double * B, double * C,
5                  double alpha, double beta,
6                  int lda, int ldb, int ldc) {
7     cilk_for (int j = 0; j < n; ++j) {
8         for (int i = 0; i < m; ++i) {
9             C[i + j * ldc] *= beta;
10        }
11    }
12    cilk_for (int jh = 0; jh < n; jh += s) {
13        int J = MIN(n-jh, s); // in case n % s != 0
14        cilk_for (int ih = 0; ih < m; ih += s) {
15            int I = MIN(m-ih, s); // in case m % s != 0
16            for (int ph = 0; ph < k; ph += s) {
17                int P = MIN(k-ph, s); // in case p % s != 0
18                for (int jl = 0; jl < J; ++jl) {
19                    int j = jh + jl;
20                    for (int pl = 0; pl < P; ++pl) {
21                        int p = ph + pl;
22                        for (int il = 0; il < I; ++il) {
23                            int i = ih + il;
24                            C[i + j * ldc] += alpha * A[i + p * lda] * B[p + j * ldb];
25                        }
26                    }
27                }
28            }
29        }
30    }
31 }
32 int main(int argc, const char * argv[]) {
33     if (argc != 4 && argc != 5) {
34         printf("Test usage: ./test m n k [s]\n");
35         exit(-1);
36     }
37     int m = atoi(argv[1]);
```

```

38     int n = atoi(argv[2]);
39     int k = atoi(argv[3]);
40     if (argc == 5) {
41         s = atoi(argv[4]);
42     }
43     // ...
44 }

```

Experiment: The best tiling step size is 128, and the running time is **0.3s(335.7088ms)**

Speedup: **6691x**, percent of peak: **8.94%**.

```

1  $ ~/opencilk/bin/clang++ gemm.cpp -fopencilk -O3 -o 6-tiling
2
3  $ ./6-tiling 4096 4096 4096 32
4  input: 4096 x 4096 x 4096
5  minimal time spent: 1455.8920 ms
6  result: correct (err = 0.000000e+00)
7
8  $ ./6-tiling 4096 4096 4096 64
9  input: 4096 x 4096 x 4096
10 minimal time spent: 708.2708 ms
11 result: correct (err = 0.000000e+00)
12
13 $ ./6-tiling 4096 4096 4096 128
14 input: 4096 x 4096 x 4096
15 minimal time spent: 335.7088 ms
16 result: correct (err = 0.000000e+00)
17
18 $ ./6-tiling 4096 4096 4096 256
19 input: 4096 x 4096 x 4096
20 minimal time spent: 554.6172 ms
21 result: correct (err = 0.000000e+00)
22
23 $ ./6-tiling 4096 4096 4096 512
24 input: 4096 x 4096 x 4096
25 minimal time spent: 1507.3377 ms
26 result: correct (err = 0.000000e+00)

```


7 Divide and Conquer

We only optimize for the case where the dimension is a power of 2 by recursively divide and conquer.

```
1  #define ISPOWER(x) ((x) & -(x)) == (x)
2  void mm_base(int m, int n, int k, double alpha,
3              const double *A, const double *B, double *C,
4              int lda, int ldb, int ldc) {
5      for (int j = 0; j < n; j++) {
6          for (int p = 0; p < k; p++) {
7              for (int i = 0; i < m; i++) {
8                  C[i + j * ldc] += alpha * A[i + p * lda] * B[p + j * ldb];
9              }
10         }
11     }
12 }
13 void mm_dac(int m, int n, int k, double a,
14             const double *A, const double *B, double *C,
15             int ldA, int ldB, int ldC) {
16     assert(ISPOWER(m) && ISPOWER(n) && ISPOWER(k));
17     if (MIN(m, n) ≤ DAC_THOLD || k ≤ DAC_THOLD) {
18         mm_base(m, n, k, alpha, A, B, C, ldA, ldB, ldC);
19     } else {
20         int& R_A = m; int& R_B = k; int& R_C = m;
21         int& C_A = k; int& C_B = n; int& C_C = n;
22         #define X(M, r, c) (M + r * (R_ ## M / 2) + c * (ld ## M) * (C_ ## M / 2))
23         cilk_spawn mm_dac(m/2, n/2, k/2, a, X(A,0,0), X(B,0,0), X(C,0,0), ldA, ldB, ldC);
24         cilk_spawn mm_dac(m/2, n/2, k/2, a, X(A,0,1), X(B,0,1), X(C,0,1), ldA, ldB, ldC);
25         cilk_spawn mm_dac(m/2, n/2, k/2, a, X(A,1,0), X(B,0,0), X(C,1,0), ldA, ldB, ldC);
26         mm_dac(m/2, n/2, k/2, a, X(A,1,0), X(B,0,1), X(C,1,1), ldA, ldB, ldC);
27         cilk_sync;
28         cilk_spawn mm_dac(m/2, n/2, k/2, a, X(A,0,1), X(B,1,0), X(C,0,0), ldA, ldB, ldC);
29         cilk_spawn mm_dac(m/2, n/2, k/2, a, X(A,0,1), X(B,1,1), X(C,0,1), ldA, ldB, ldC);
30         cilk_spawn mm_dac(m/2, n/2, k/2, a, X(A,1,1), X(B,1,0), X(C,1,0), ldA, ldB, ldC);
31         mm_dac(m/2, n/2, k/2, a, X(A,1,1), X(B,1,1), X(C,1,1), ldA, ldB, ldC);
32         cilk_sync;
33     }
34 }
```

However, the result is that the program cannot suck all CPU power due to the synchronization, and the best result is **521.0503ms** for `DAC_THOLD=128`.

8 Compiler Flags

```
1 $ ~/opencilk/bin/clang++ gemm.cpp -fopencilk -O3 -o 8-avx -march=native -ffast-math
2 $ ./8-avx 4096 4096 4096
3 input: 4096 x 4096 x 4096
4 minimal time spent: 323.6070 ms
5 result: correct (err = 4.973799e-14)
```

Fast math doesn't guarantee the `err` to be 0, but it is faster.

The running time is **0.32s**.

Speedup: **6941x**, percent of peak: **9.27%**.